



The easy and lightweight ORM for Python

Emmanuel Turlay
Instacart
emmanuel@turlay.net

What's an ORM?

Object Relational Mapping

Classes



Database



SQLAlchemy, ActiveRecord, DataMapper...

Why an ORM?

```

114 def save_to_db(x, type, apply_actuals=False):
115     value = x.feasible_levels
116     hour = "%s:00" % str(int(x.hour))
117     shift_type = x.shift_type
118
119     if shift_type == "picking_only":
120         params = (datetime.utcnow(), int(x.zone_id), x.date, hour, int(x.warehouse_id), int(x.warehouse_location_id), shift_type)
121         if type == 'update':
122             params = (value, value,) + params
123             sql = """UPDATE staffing_levels SET num_shoppers = %s, alternate = %s, updated_at = %s WHERE zone_id = %s AND date=%s AND local_start_time=%s::time AND warehouse_id =
                %s AND warehouse_location_id = %s AND shift_type = %s"""
124
125         elif type == 'insert':
126             params = (value, value, datetime.utcnow()) + params
127             sql = """INSERT INTO staffing_levels
128                 (num_shoppers, alternate, duration_in_minutes, created_at, updated_at, zone_id, date, local_start_time, warehouse_id, warehouse_location_id, shift_type)
129                 VALUES (%s, %s, 60, %s, %s, %s, %s, %s, %s, %s, %s)"""
130
131         # log.info("%s zone_id %s, wlid %s, date %s, hour %s, value %s, shift_type %s" % (type, x.zone_id, int(x.warehouse_location_id), x.date, hour, value, shift_type,))
132         execute(sql, params)
133
134     else:
135         params = (datetime.utcnow(), int(x.zone_id), x.date, hour, shift_type)
136         if type == 'update':
137             if apply_actuals:
138                 params = (value,) + params
139                 query_part = """SET num_shoppers = %s, """
140             else:
141                 params = (value, value,) + params
142                 query_part = """SET num_shoppers = %s, alternate = %s, """
143             sql = """UPDATE staffing_levels """ + query_part + """updated_at = %s WHERE zone_id = %s AND date=%s AND local_start_time=%s::time AND shift_type = %s"""
144
145         elif type == 'insert':
146             params = (value, value, datetime.utcnow(),) + params
147             sql = """INSERT INTO staffing_levels
148                 (num_shoppers, alternate, duration_in_minutes, created_at, updated_at, zone_id, date, local_start_time, shift_type)
149                 VALUES (%s, %s, 60, %s, %s, %s, %s, %s, %s)"""
150
151         # log.info("%s zone_id %s, date %s, hour %s, value %s, shift_type %s" % (type, x.zone_id, x.date, hour, value, shift_type,))
152         execute(sql, params)

```

Abstraction layer

Database connections

Interact with database objects

Error handling

Model-level business logic




jardin (*noun, french*) – garden, yard, grove



pandas

$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$



Configure

```
# jardin_conf.py

import logging

DATABASES = {
    'db1_master': 'postgres://user:pass@host:port/database_w',
    'db1_replica': 'postgres://user:pass@host:port/database_r',
    'db2': {
        'scheme': 'mysql',
        'username': 'user',
        ...
    }
}

LOG_LEVEL = logging.DEBUG

WATERMARK = 'MyCoolApp'
```

```
$ export JARDIN_CONF=/path/to/jardin_conf.py
```


Declare

```
# db_models.py

import jardin

class Db1AbstractModel(jardin.Model):
    db_names = {
        'master': 'db1_master',
        'replica': 'db1_replica'
    }

class Db2AbstractModel(jardin.Model):
    db_names = {
        'master': 'db2',
        'replica': 'db2'
    }

class User(Db1AbstractModel): pass

class Order(Db2AbstractModel): pass
```

```
>>> from db_models import User

>>> User.insert(values={'name': 'jardin'})
DEBUG:jardin:INSERT INTO users VALUES (...) /* MyCoolApp */
User(id=2, name='jardin', created_at='...', updated_at='...', ...)

>>> User.select(where={'active': False})
DEBUG:jardin:SELECT * FROM users WHERE ... /* MyCoolApp */
  id  name  active
0   2  'jardin'  False
1   1  'sqlalchemy' False

>>> User.update(values={'active': True}, where={'name': 'jardin'})
DEBUG:jardin:UPDATE users SET ... WHERE ... /* MyCoolApp */

>>> user = User.find(1)
DEBUG:jardin:SELECT * FROM users WHERE ... /* MyCoolApp */

>>> user.active = False

>>> user.save()
DEBUG:jardin:UPDATE users SET ... WHERE ... /* MyCoolApp */

>>> User.delete(where={'active': False})
DEBUG:jardin:DELETE FROM users WHERE ... /* MyCoolApp */
```

```

# path/to/file.py

from datetime import datetime, timedelta
from db_models import User, UserSettings
from jardin.comparators import gt

users = User.select(
    select={
        'user_id': 'id',
        'username': 'name',
        'newsletter': 'us.newsletter'
    },
    where={
        'active': True,
        'created_at': gt(datetime.utcnow() - timedelta(days=30)),
        'deleted_at': None
    },
    inner_join=[UserSettings]
    order={'created_at': 'DESC'},
    limit=10
)

```

```

DEBUG:jardin:('SELECT id AS user_id, name AS username, us.newsletter AS
newsletter FROM users u INNER JOIN user_settings us ON us.user_id =
u.id WHERE active IS TRUE AND created_at > %(created_at)s AND
deleted_at IS NULL ORDER BY created_at DESC LIMIT 10; /* MyCoolApp | /
path/to/file.py:8 */', OrderedDict(('created_at', '2018-04-09
19:00:00'))))

```

```
import jardin
```

```
dataframe = jardin.query(  
    sql='SELECT * FROM orders WHERE abc = %(abc)s LIMIT 10;'  
    params={'abc': 123},  
    db='db1_replica'  
)
```

```
dataframe = jardin.query(  
    filename='/path/to/extract.sql'  
    params={'abc': 123},  
    db='db1_replica'  
)
```

```
import pandas as pd
from db_models import User

df = pd.DataFrame(...)

with User.transaction():
    User.delete(where={'id': df.id})
    User.insert(values=df)
```

Python 2.7+ and 3.5+

Pandas-dataframe integration

PostgreSQL, MySQL, SQLite, AWS Redshift and Snowflake

Multiple databases with master/replica split

Transactions

Connection drop recovery

Single-statement dataframe insertion

À la ActiveRecord query scopes

created_at, updated_at and soft-deletes out-of-the-box support

pip install jardin

instacart.github.io/jardin